

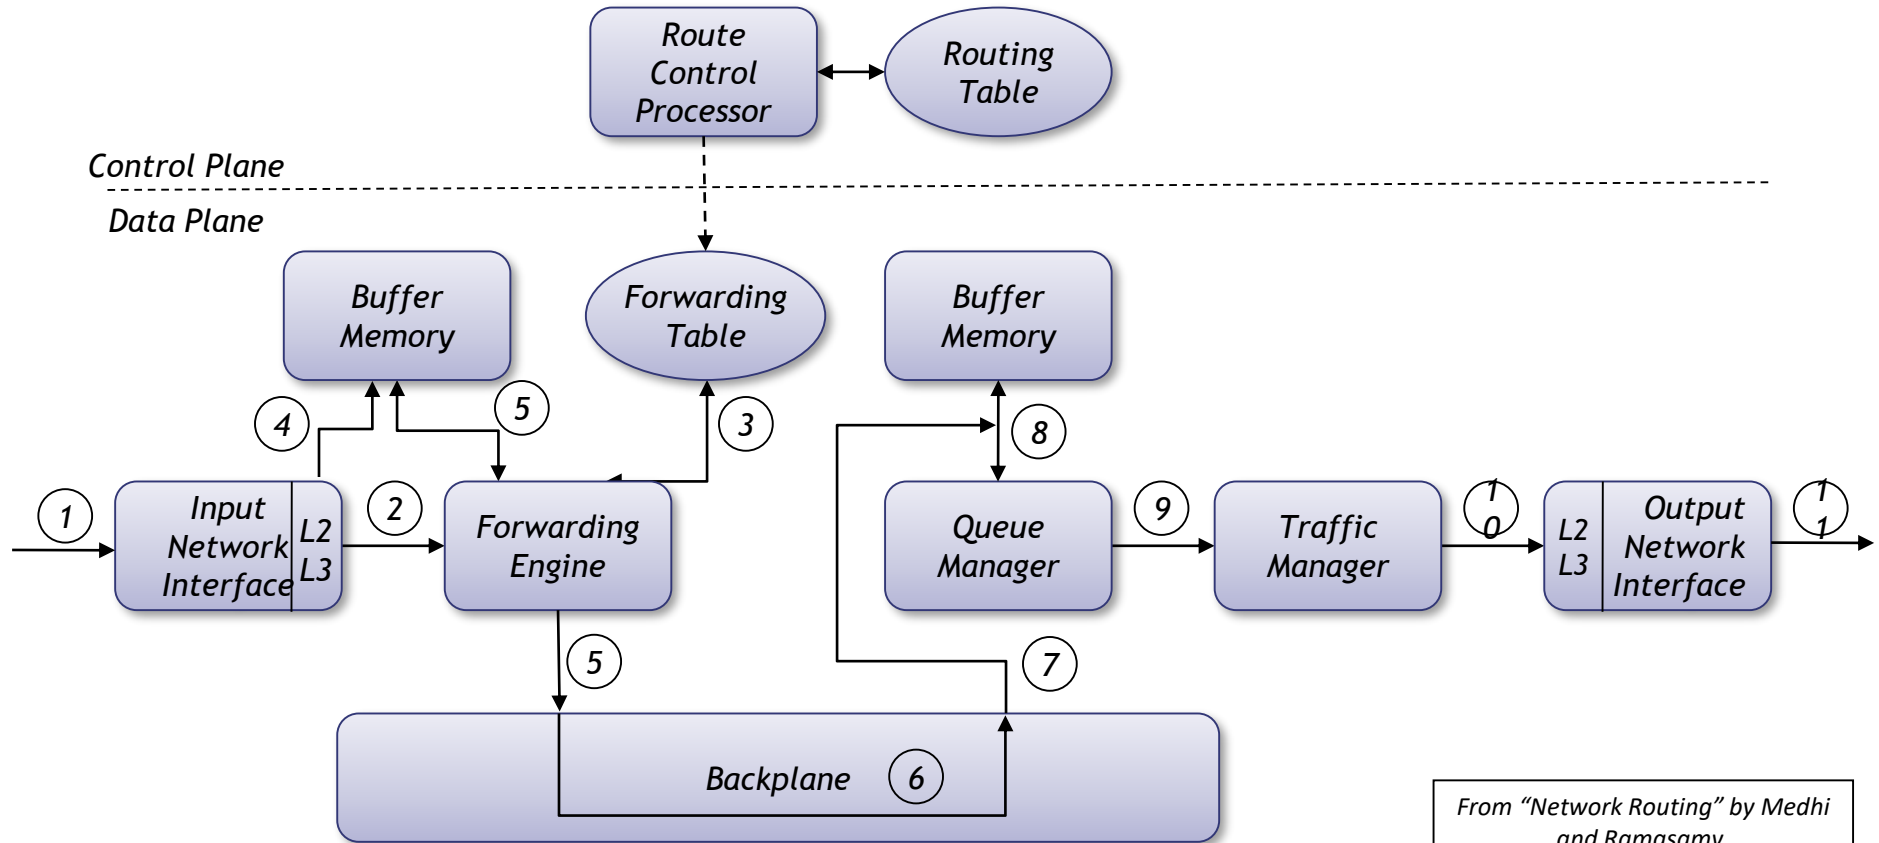
# Packet-Switching

## Switching Inside Routers

# Outline

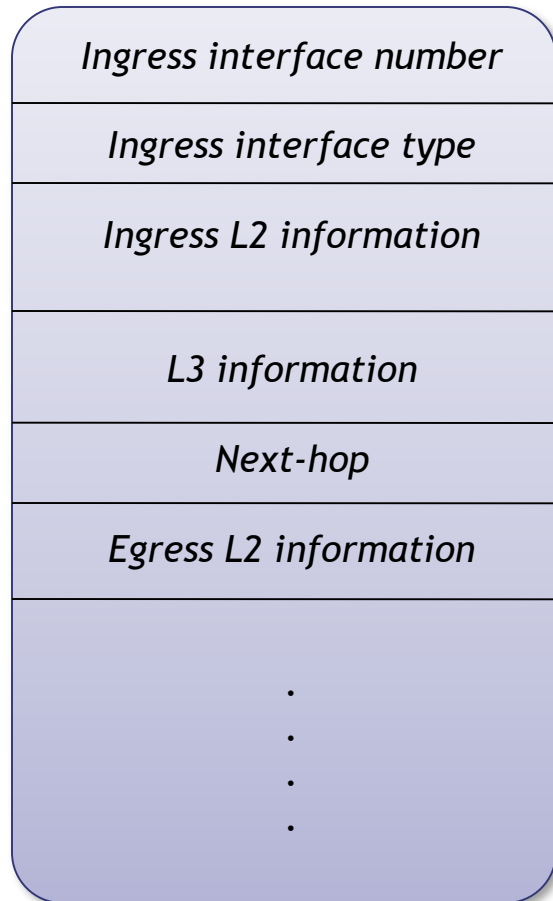
- **Packet flow in a router**
- Fast path vs. slow path
- Memory organization for packet switching
- Hardware-aided switching
- Switch fabric

# Packet Flow



From "Network Routing" by Medhi and Ramasamy.  
Morgan Kaufmann Publishers

# The packet context



- A data structure (an “internal packet”) that essentially serves as a scratch pad for carrying information between different stages of packet processing inside the router.

*From “Network Routing” by Medhi and Ramasamy.  
Morgan Kaufmann Publishers*

# Ingress packet processing (1)

- 1) An IP packet is received at the Input Network Interface (assume it is an Ethernet packet, so the network interface is an Ethernet port). Then the interface:
  - a) Interprets the Ethernet header,
  - b) Detects frame boundaries,
  - c) Identifies the starting point of the IP packet,
  - d) The L2 processing logic removes the L2 header and appends it to the *packet context* (i.e., source and destination MAC addresses),
  - e) The L2 processing logic takes the payload (the IP packet) and along with the packet context sends it to the L3 processing logic,

# Ingress packet processing (2)

- f) The L3 processing logic locates the IP header and checks its validity,
  - g) The L3 processing logic extracts the relevant IP header information (DA, SA, protocol type, DSCP bits for differentiated services, type of load – TCP or UDP -, destination and source ports) and stores it in the packet context.
- 2) The packet context is sent to the Forwarding Engine.
  - 3) The Forwarding Engine searches the Forwarding Table to determine the next-hop. The next-hop information contains the egress line card and the outgoing port the packet needs to be transferred to. This information is populated in the packet context.

# Ingress packet processing (3)

- 4) While the Forwarding Engine is determining the next-hop using the packet context, the L3 processing logic sends the IP packet to be stored in the Buffer Memory temporarily.
- 5) When the Forwarding Engine finishes, the packet context is appended to the address of the packet in memory and sent to the Backplane Interface.
- 6) From the packet context, the backplane knows to which line card the packet needs to be transferred. It extracts it from memory and together with the packet context, sends them to the appropriate egress card. Priority of the packet (as decided by the Forwarding Engine) is taken into account while transmitting on the backplane.

# Egress Packet Processing (1)

- 7) The backplane interface on the egress line card receives the packet and stores it in the line card memory.**
- 8) The received packet context is updated with the new address of the memory location and sent to the Queue Manager.**
  - a) The queue manager examines the packet context to determine the packet priority (the priority was determined by the forwarding engine in the ingress line card during packet classification).**
  - b) Next the queue manager inserts the context of the packet in the appropriate queue.**



# Egress Packet Processing (2)

- c) The queue manager implements a scheduling algorithm. The scheduling algorithm chooses the next packet to be transmitted. In some instances, the queues could be full because of congestion in the network. In order to handle such cases, the queue manager implements packet dropping behavior to proactively drop packets when the router experiences congestion.
- 9) Once the packet is scheduled to be transmitted, the Traffic Manager examines its context to identify the customer and if there are any transmit rate limitations that need to be enforced according to the service contract (traffic shaping). If the traffic exceeds any rate limitations, the traffic manager delays or drops the packet in order to comply with the agreed rate.

# Egress Packet Processing (3)

- 10) Finally, the packet context arrives at the egress network interface. The packet is retrieved from memory. L3 processing logic updates its TTL and updates the checksum. The L2 processing logic adds the appropriate L2 headers.
- 11) The packet is transmitted.

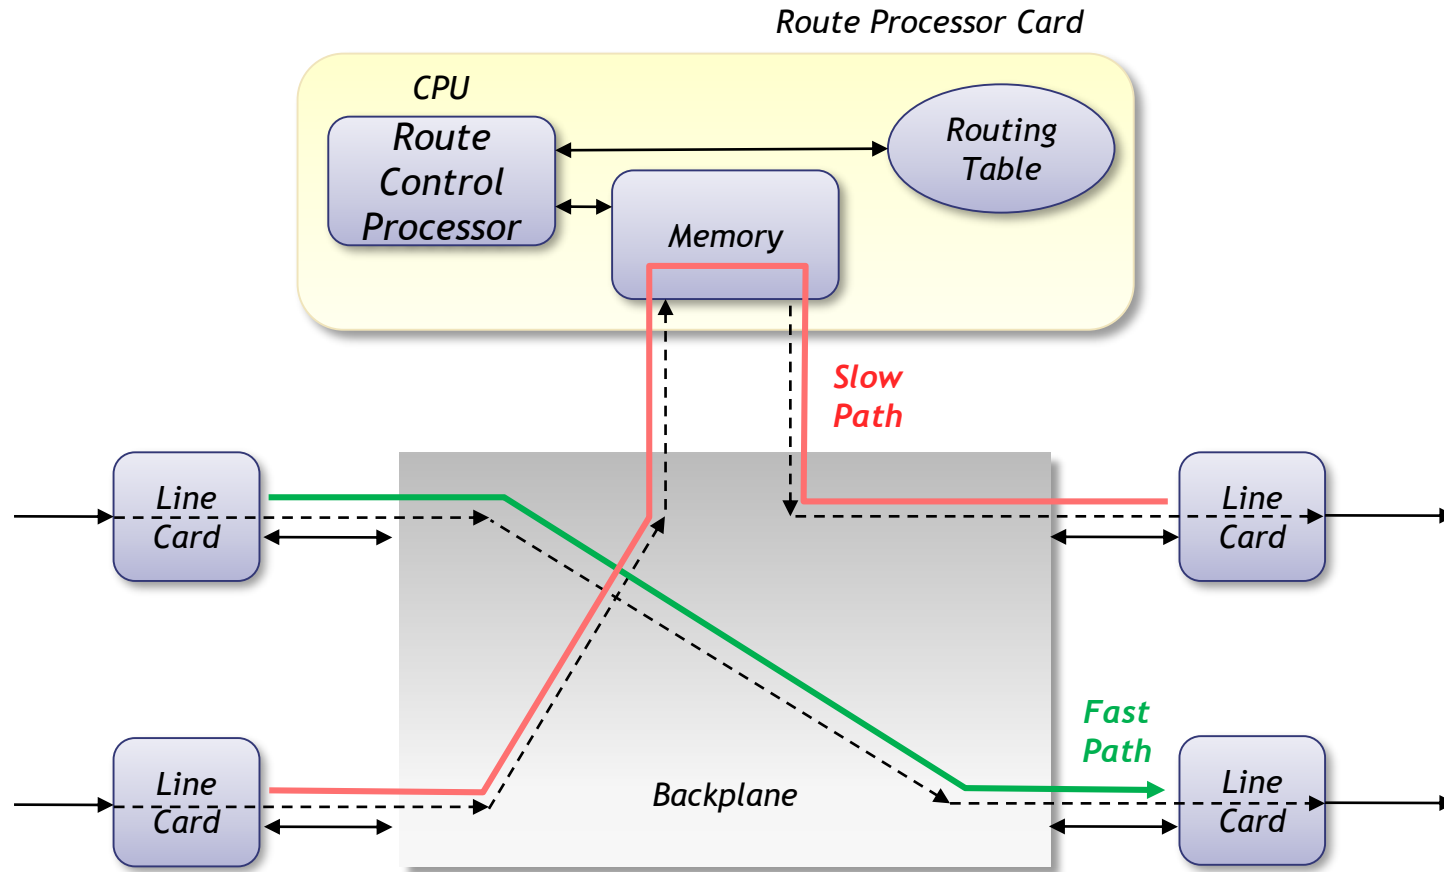
# Index

- Packet flow in a router
- **Fast path vs. slow path**
- Memory organization for packet switching
- Hardware-aided switching
- Switch fabric

# Packet processing: Fast vs. Slow Path

- **The tasks performed by a router can be categorized into:**
  - ❖ **Time-critical (fast path):** deals with the majority of the packets and needs to be highly optimized to achieve gigabit/s forwarding rates. Implemented in hardware in high performance routers. Can be broadly grouped into:
    - ✓ Header processing: packet validation, packet lifetime control, TTL and checksum calculation.
    - ✓ Forwarding: destination address lookup, packet classification for service differentiation, packet buffering, queue management and scheduling.
  - ❖ **Non-time-critical (slow path):** Management, control and error handling. Normally implemented in software (CPU). Examples:
    - ✓ Processing of data packets that led to errors in the fast path and generation of ICMP packets to inform the originating source of the packets
    - ✓ Processing of routing protocol keep-alive/update messages from/to adjacent neighbors
    - ✓ Processing/generation of packets pertaining to management protocols, such as SNMP
    - ✓ ARP
    - ✓ Packet Segmentation
    - ✓ IP header options: timestamp, source routing, route record, router alert

# Packet processing: Fast vs. Slow Paths



From "Network Routing" by Medhi and Ramasamy.  
Morgan Kaufmann Publishers

# Index

- Packet flow in a router
- Fast path vs. slow path
- **Memory organization for packet switching**
- Hardware-aided switching
- Switch fabric

# Memory Organization: Pools and particle buffers

- Significant amount of memory used to queue packets
- Optimization: organize in pools
  - ❖ A pool is a large array of memory blocks of the same size (derived from the MTU)
    - Problem: Packets are normally smaller than the link's MTU
    - Solution: A number of pools of different sizes
- Alternative: Divide memory in small blocks: Particles dividing the packet in several particles.

# Memory Organization: Public vs. Private Buffers

- **Problem: If all buffers shared:**
  - ❖ One interface can take all the capacity
    - ✓ Starvation for the rest of the interfaces.
- **Solution:**
  - ❖ Some buffers are assigned for exclusive use of each interface (private buffers)
  - ❖ The remaining public buffers / system buffers:  
Shared among all interfaces



# Index

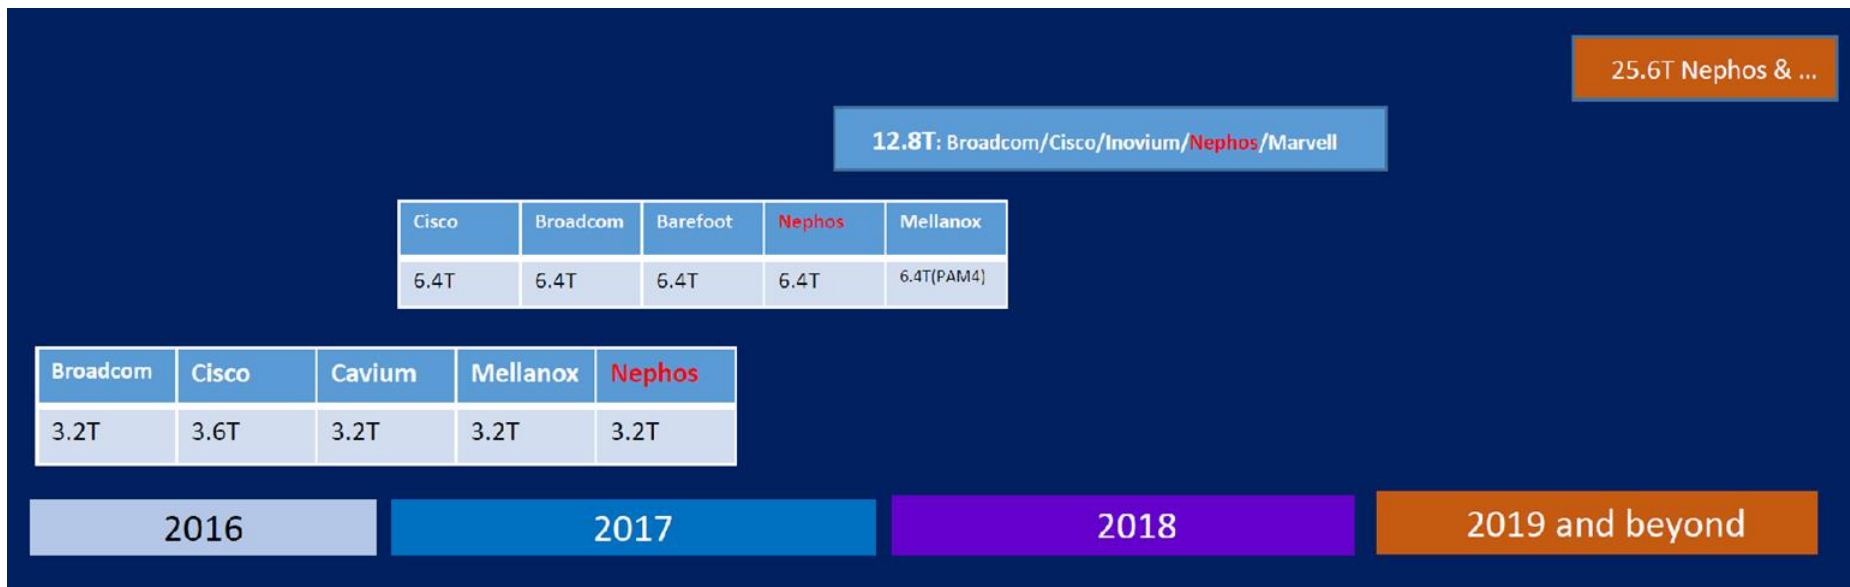
- Packet flow in a router
- Fast path vs. slow path
- Memory organization for packet switching
- **Hardware-aided switching**
  - ❖ ASICs
  - ❖ Network processors
  - ❖ TCAMs/search engines
  - ❖ FPGAs
- Switch fabric

# Switching Hardware

- All packet processing can be done in software
- But: Parts of the process can be realized in hardware to speed up the process, in particular for high performance routers.
- Examples
  - ❖ Custom ASICs for routers (Cisco, Broadcom, Mellanox,...)
  - ❖ Network Processors (NP) (Broadcom, Mellanox,...)
  - ❖ Ternary Content Addressable Memories (TCAMs) /Search Processors (Renesas, Broadcom,...)
  - ❖ FPGAs (Xilinx, Intel,....)

# Switching Hardware: ASICs

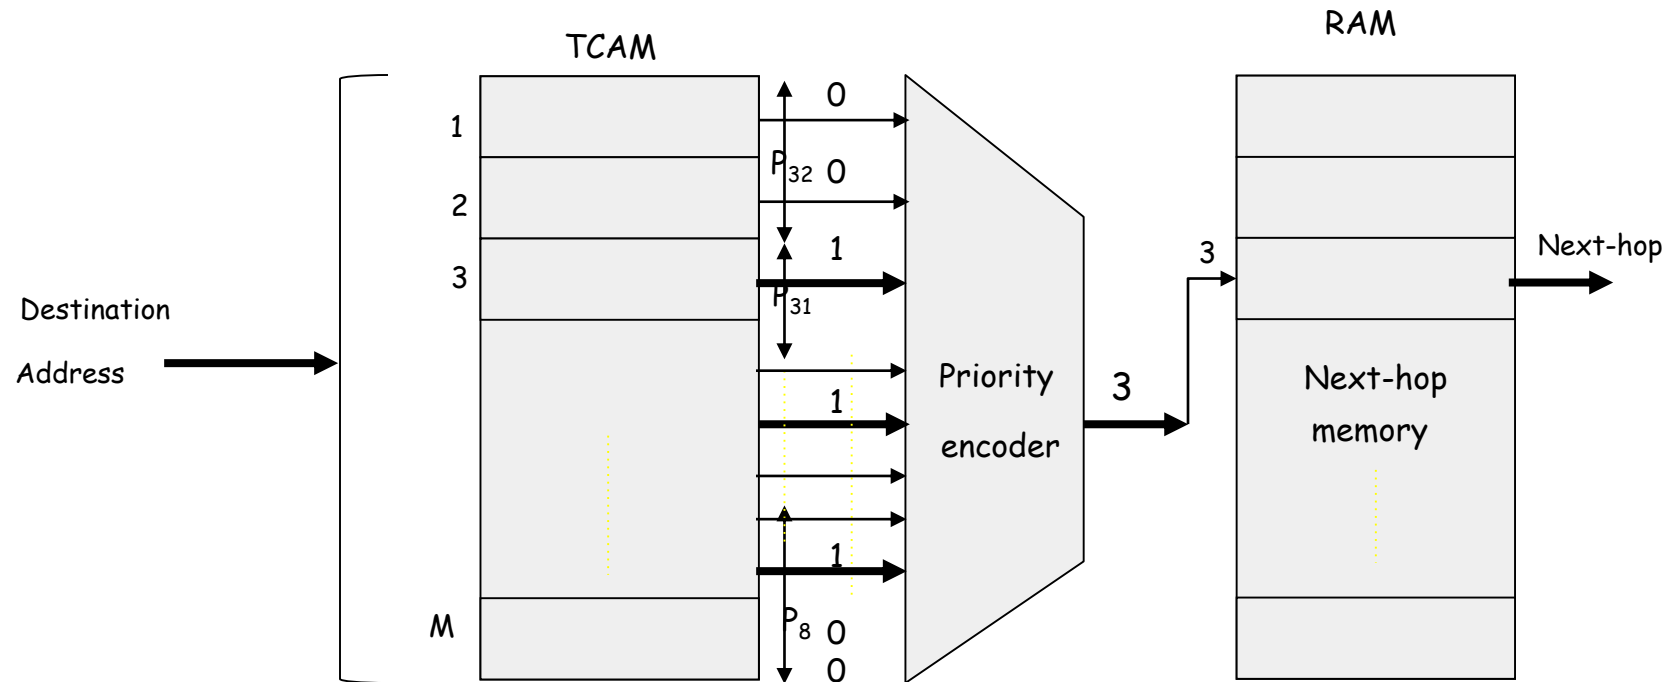
- Higher performance
- Harder to make them programmable
- High development cost (and time)
- System on a Chip (SoC)



# Switching Hardware: Network processors

- **Can we have a technology that :**
  - ❖ Has the flexibility of programmable processors ?
  - ❖ Has the high speed of ASICs ?
- **Network processor: programmable as a CPU, performance close to ASIC (due to multi-processors)**
- **Network Processor:**
  - ❖ Specialized circuitry and micro-engines to perform all generic packet processing functions.
  - ❖ A major programmable module, usually a tailored-made RISC CPU.
  - ❖ Real time operating system

# Switching Hardware: TCAMs (route lookup)



- **Ups**

- ❖ High capacity: e.g. 128 K entries each 72-bits.
- ❖ High performance > 100 million lookups per sec.

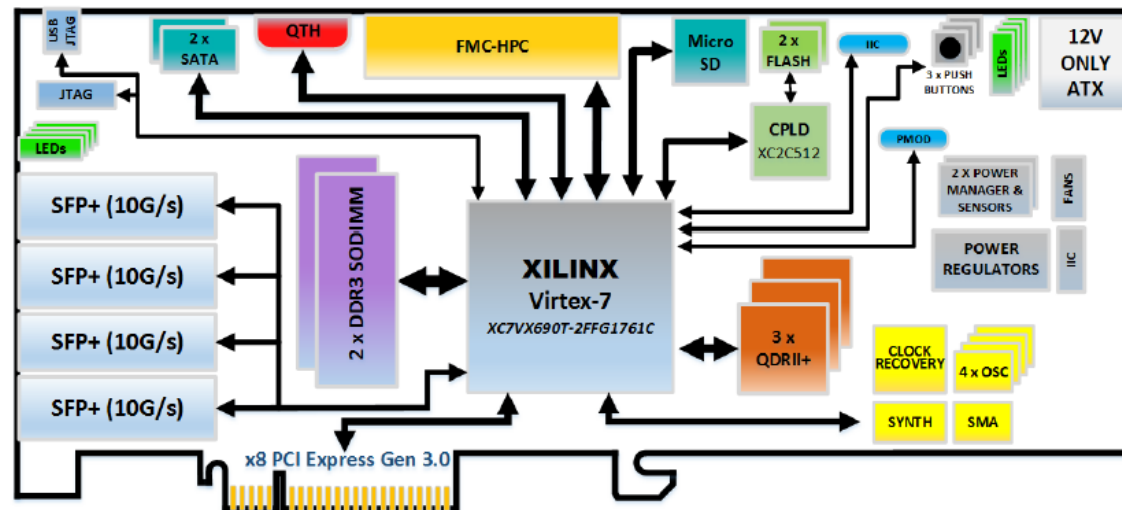
- **Cons**

- ❖ High cost
- ❖ High power consumption

# Switching Hardware: FPGAs

- Field Programmable Gate Arrays
- Have evolved to be large SoC with embedded processors, logic blocks, memories and interfaces
- Can be used to implement reconfigurable network devices
- Flexibility and performance
- Example NetFPGA: N. Zilberman et al, "NetFPGA SUME: Toward 100 Gbps as Research Commodity," IEEE Micro, 2014.

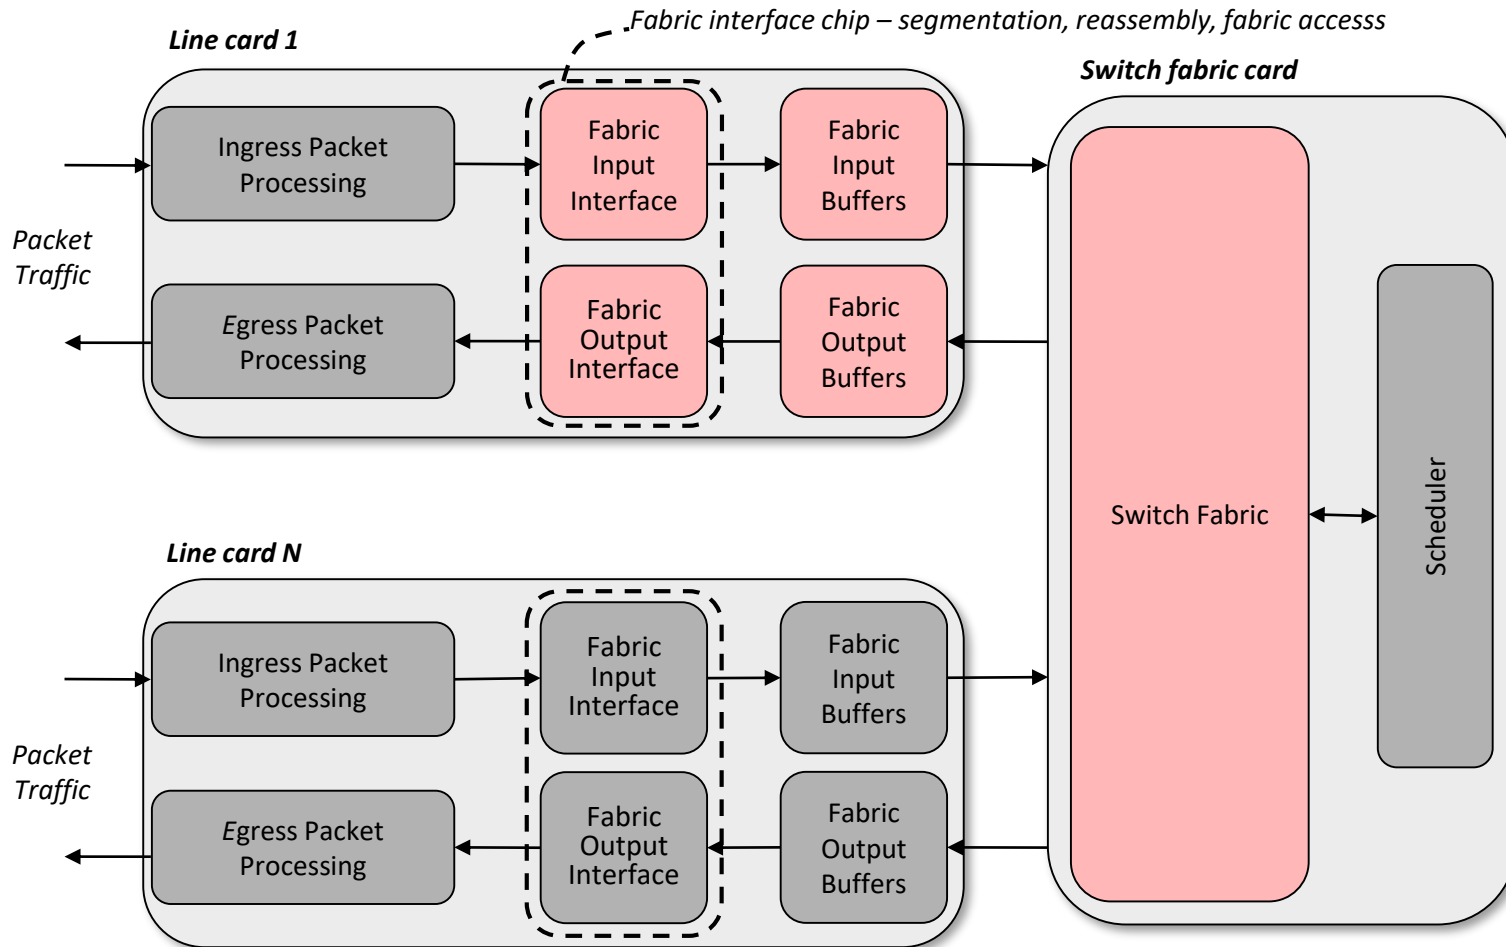
[http://ieeexplore.ieee.org/xpls/abs\\_all.jsp?arnumber=6866035](http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=6866035)



# Index

- Packet flow in a router
- Fast path vs. slow path
- Memory organization for packet switching
- Hardware-aided switching
- **Switch fabrics**
  - Switch fabric design
  - HOL → VOQ problem in input-buffered crossbar switch fabric
  - Control algorithms

# Generic Switch Architecture






# Switching Fabric Design

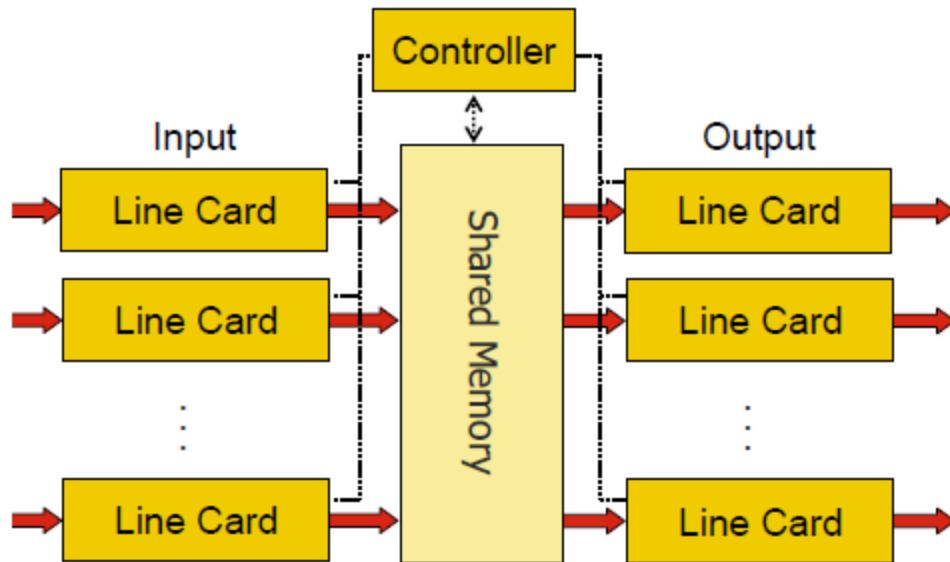
- **It is important to control the switching fabric**
  - ❖ Calculate and establish the connections
- **Problem: Variable packet sizes**
  - ❖ Need to calculate connections taking into account that each cross-point will be available at a different time (when the packet finishes transmission through the fabric)
  - ❖ High complexity of control
- **Thus: some fabrics have a fixed-size switching unit**
  - ❖ Packets are split into chunks or cells.
  - ❖ Decisions are taken at fixed intervals
  - ❖ Much simpler scheduling
  - ❖ But: Additional circuitry to fragment and reassemble packets (SAR)

# Switch Fabrics: Shared Medium (Bus)



- Operation:
  - Time-sharing of medium (TDM)
  - Each output monitors address tag to identify its packets
- Physical limit of speed of bus ( $N \times R$ ), address filters and output buffers
  - Address filters and output buffers must operate at  $N \times R$  ( $N$ = number of ports)
- If bus speed  $\geq N \times R \Rightarrow$  Output queuing: optimal throughput
- No sharing of buffers among line cards
  - more memory required for the same packet loss w.r.t. a shared memory fabric
- Natural implementation of multicast and broadcast


# Switch Fabrics: Shared Memory



- Operation
  - Packets or cells stored in multiple shared memory banks (possibly interleaved access)
- Kind of output queuing approach where all output buffers belong to the shared memory pool
- Popular switch architecture (efficient memory usage with unbalanced traffic)

- **Physical limit of access time of RAM (read/write at speed  $\underline{2N}$  times the speed of ports. i.e.  $2NR$ )**
- **Processing by central controller at packet rate**
  - not trivial if the controller has to handle multiple priority classes and complicated packet scheduling
- **Sharing of buffers (output buffers)**
  - Efficient use of memory due to statistical sharing. Minimum memory for a target loss rate.
- **Multicast requires extra circuitry (multiple reads, or multiple copies in memory)**
- **E.g. Catalyst 8510, Juniper M40**

# Switch Fabrics: Crossbar Switch: Input-Queued Crossbar Switch (IQ)



- Non-blocking switch
- Number of cross points  $N \times N$
- Memory grows lineally with  $N$ 
  - E.g. Cisco 12000

- **Input buffering**

- **Ups:**


- Crossbar and memory run at line rate
- Highly scalable architecture

- **Cons:**

- ❖ Head-of-line (HOL) blocking: 58.6% w/o arbiter with uniform traffic

# Switching Fabric Control: Head of Line Blocking


- Problem in crossbar switch: Two packets compete for the same outgoing interface
  - ❖ They block the following queued packets on that interface (input queuing) □ Head of Line (HOL) blocking



# Switching Fabric Control: VOQ (Anderson 1993)


- **Solution to avoid HOL Blocking:**

- ❖ **Virtual Output Queuing (VOQ):** a separate queue for each outgoing interface at the input buffer



# Switch Fabrics: Crossbar Switch II: Output-Queued Crossbar Switch

N X N Crossbar Switch



- Non-blocking switch
- Number of cross points  $N \times N$
- Memory grows linearly with  $N$
- **Output Queued**
- **Ups:**
  - No Head-of-line (HOL) blocking: 100% performance
  - No input traffic scheduling is required !
- **Cons:**
  - Crossbar and memory need to run at  $R \times N$  (speed-up =  $N$ )
  - Less scalable architecture

# Combined Input-Output Queued crossbar (CIOQ)

- Why do Input-buffered switch fabrics add delay?


- ❖ Output contention:
  - ✓ Several inputs need to transmit to the same output.
- ❖ Input queuing:
  - ✓ HOL blocking: a packet waiting for an output to get free blocks the packets behind it, even if their output ports are available. -- solvable: VoQ
  - ✓ Queueing: Assumed VoQ...: two packets on the same input go to different free outputs. One must wait.

- Solutions:

- ❖ VoQ with appropriate scheduling solves HOL factor
- ❖ Increase the switch fabric speed to  $S$  times faster than the line rate. This is called "Combined Input-Output Queued" crossbar (\*)
  - ✓ Theoretically equal to the number of ports:  $S=N$ .
  - ✓ Usually  $S=2$  to 5
- ❖ Assign priorities to packets.
  - ✓ Packets with higher priority will experience less delay.


(\*) A speedup requires buffering at the output

N X N Crossbar Switch





# Switch Fabrics: Distributed Output-Buffered Switch



- **Immediate placement of packets in output buffers**
- **Ups**
  - Switch fabric (N-buses), AF and buffers run at line-speed  $R$  - no need to speed-up
  - Optimal throughput: 100%
  - Better latency control (QoS guarantees, switching time does not depend on arbiter algorithms)
- **Cons**
  - Cost: Quadratic  $N^2$  growth of buffers
  - Modularity: limit to number of ports  $N$
- **Lower cost Alternative:**
  - Knock-out switch: add a concentration stage to access to  $L < N$  buffers.  $L=8$  for loss= $10^{-6}$  and  $N$  large

# Conclusions on switch fabric design

## ● High-end Packet switches (packet/cell)

- ❖ Initially shared-memory or output-queued (ATM)
- ❖ Were difficult to scale to high throughputs
  - ✓ Memory access time has not been reduced for the last ten years: real bottleneck of switches
- ❖ Conclusion:
  - ✓ DOQ (Distributed output-buffered) provides 100% throughput and delay guarantees but high memory requirements
  - ✓ Try to emulate DOQ with cheaper approaches:
    - CIOQ requires some speed-up and two queueings
    - IQ switches seem to be more scalable as they require fewer buffers, but they require scheduling to get near 100% of throughput
      - So let's find simple scheduling algorithms...

# Switching Fabric Control Algorithms

- **Requirements:**

- ❖ Simple to implement
- ❖ Fast
- ❖ Fair
- ❖ Starvation free ( no endless waiting in VOQ )
- ❖ High throughput


- **Algorithms:**

- ❖ Parallel Iterative Matching (PIM)
- ❖ Round Robin Matching (RRM)
- ❖ iSLIP

# Parallel Iterative Matching (PIM)

- **Three steps in one iteration:**
  1. **Request:** Each input sends a request to every output for which it has a queued cell.
  2. **Grant:** If an output receives several requests it chooses one randomly and notifies each input whether its request was granted or not.
  3. **Accept:** If an input receives several grants it chooses one randomly.
- **Repeat steps 1-3:**
  - ❖ Only with free inputs and outputs !!
  - ❖ When no new matching can be found, the algorithm stops.

# PIM: Example (Two Iterations)



● Free / Unmatched

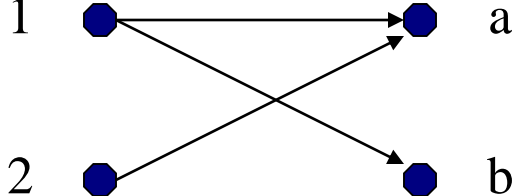
# PIM: Discussion

- **Several iterations necessary**

- ❖ On average 75% matches per iteration


- **Problem:**

- ❖ Generation of random decisions at high speed is not trivial (three-step handshake)

- ❖ Probably low fairness. Example:
- 
- The diagram shows two input nodes on the left, labeled 1 and 2, and two output nodes on the right, labeled a and b. Node 1 is connected to both node a and node b. Node 2 is also connected to both node a and node b. This represents a fully connected bipartite graph where each input has a probability of 1/2 of being matched to either output.

- ❖ Output *a* will grant request of each input with probability  $\frac{1}{2}$
- ❖ Input 1 will only accept the grants  $\frac{1}{2}$  of the times.
- ❖ So for fully loaded inputs (100 %) the throughput is:
  - ✓  $T_{1,a} = 1/4$ ;  $T_{1,b} = 3/4$ ;  $T_{2,a} = 3/4$

# PIM: Discussion




$$P_{1,a} = 1/4; P_{1,b} = 3/4; P_{2,a} = 3/4$$

# RRM – Round-Robin Matching


- RRM chooses following a round-robin scheme instead of randomly. Each “selector” keeps a pointer to the highest priority port,  $a_i$  (accept pointer) and  $g_j$  (grant pointer). Each iteration has 3 steps, in which the inputs and outputs not assigned are used.
- The algorithm finishes when in an iteration there are no new connections.
  - Request: Each input sends a request to every output for which it has a queued cell.
  - Grant: If an output receives several requests, it chooses the first one starting by the highest priority element, marked by  $g_j$ . The pointer is updated to the value of the next input to the granted one.
  - Accept: If an input receives several grants, it chooses the first one starting by the highest priority element, marked by  $a_i$ . The pointer is updated to the value of the next output to the granted one.




# RRM: Example I



Step 1: Request




Step 2: Grant




Step 3: Accept


# RRM: Example II



Step 1: Request



Step 2: Grant




Step 3: Accept

# RRM: Starvation

- **Problem: RRM is not starvation free**

- ❖ **Example:**



- ❖ **Assumptions:**


- ✓ There are always cells waiting to be transferred.
- ✓ The destination is always the same.

- **Sequence of grants / accepts**


- ❖ Outputs 1 and 3: always grant input 1
- ❖ Output 2: always grants input 1 at first (updates pointer  $g_2$ !)
- ❖ But: input 1 always ignores this grant (and data never goes to output 2)  
Instead it accepts output 1 indefinitely.
- ❖ Output 2 can grant either input 2 or input 3 only.
- ❖ Cells from input 1 to output 2 will never be granted.

# RRM: Starvation Example I


## Step 1 – 1st Iteration



Step 1: Request




Step 2: Grant




Step 3: Accept

# RRM: Starvation Example I

## 1<sup>st</sup> Step – 2nd Iteration




Step 1: Request



Step 2: Grant




● Without assignment



Step 3: Accept

# RRM: Starvation Example II

## 2<sup>nd</sup> Step – 1st Iteration




# iSLIP (McKeown 1995)


- $g_j$  are updated only if their grants are accepted. Starvation is avoided since the selected pair is given the minimum priority. Each iteration has 4 steps, in which the not assigned inputs and outputs are used. The algorithm finishes when in an iteration there are no new connections.
  - ❖ Request: Each input sends a request to every output for which it has a queued cell.
  - ❖ Grant: If an output receives several requests, it chooses the first one starting by the element of highest priority, marked by the  $g_j$ .
  - ❖ Accept: If an input receives several grants, it accepts the first one starting by the element of highest priority, marked by  $a_i$ . The pointer is updated to the value of the next output to the accepted one. It is updated only in the first iteration.
  - ❖ Update: The  $g_j$  that have been accepted by an input are updated. The  $g_j$  are updated only in the first iteration.

# iSLIP: Example


## 1st Step – 1st Iteration



Step 1: Request



Step 2: Grant





Step 3+4: Accept+Update




# iSLIP: Example

## 1st Step – 2nd Iteration




Without assignment




Step 3+4: Accept+Update

# iSLIP: Example


## 2nd Step – 1st Iteration



Step 1: Request



Step 3+4: Accept+Update



Step 2: Grant

(...Further iterations possible)

□ Now, the system has not arrived in the same state!

# Performance Comparison




Fig. 5. Performance of RRM and SLIP compared with PIM for i.i.d. Bernoulli arrivals with destinations uniformly distributed over all outputs. Results obtained using simulation for a  $16 \times 16$  switch. The graph shows the average delay per cell, measured in cell times, between arriving at the input buffers and departing from the switch.

# References

- R. Rojas-Cessa, "Interconnections for Computer Communications and Packet Networks", ISBN 9781482226966, CRC Press, 2017.
- R. White, V. Bollapragada, and C. Murphy, "Inside Cisco IOS Software Architecture". Cisco Press 2000 (Available in the library).
  - ❖ Chapters 2, 3, 5
- H.J. Chao, C.H. Lam, E.Oki, "Broadband Packet Switching Technologies", Wiley & Sons, 2001
  - ❖ Chapter 1—3
- H. J. Chao, "Next Generation Routers" in: Proceedings of the IEEE September 2002. ([www.ieee.org](http://www.ieee.org))
- F. Baker, "Requirements for IPv4 Routers", RFC 1812, IETF ([www.ietf.org](http://www.ietf.org))
- C. Semeria, "Internet Backbone Routers and Evolving Internet Design", Juniper Networks 2000. ([www.juniper.net](http://www.juniper.net))
- C. Semeria, "Implementing a Flexible Hardware-based router for the new IP infrastructure", Juniper Networks 2000. ([www.juniper.net](http://www.juniper.net))